

Stroke_analysis_notebook

January 21, 2025

1 Stroke Dataset Analysis

Stroke is a serious medical condition that occurs when the brain's blood supply is suddenly interrupted, causing damage to brain tissue. It is one of the leading causes of disability and death worldwide, significantly affecting individuals, families, and healthcare systems. Detecting stroke risks early and taking timely action are essential to improving patient outcomes and lessening its overall impact.

This project uses machine learning to develop a predictive model that can effectively identify individuals at risk of stroke. This dataset contains information from 5,110 people and captures various demographic details such as age, gender, marital status, and type of work and health-related factors like hypertension, heart disease, average glucose levels, BMI, and smoking habits.

Our goal is to create a classification model that predicts the likelihood of stroke, where the target variable is binary (0 indicating no stroke and 1 indicating stroke). By analyzing these variables and applying machine learning techniques, the model aims to help healthcare professionals identify high-risk individuals. This can enable earlier interventions and preventive care, ultimately contributing to better health outcomes and more effective stroke management. The purpose of this project is to create classification models that predicts the likelihood of the stroke using a binary target variable(0 indicates no stroke, 1 indicates stroke)

In this project, we will:

- 1.Explore the dataset to understand its structure and contents.
- 2.Preprocess the data to address missing values and convert categorical variables into usable formats.
- 3.Engineer features to extract meaningful insights.
- 4.Experiment with different machine learning algorithms(gradient boosting, random forest,logistic regression, kmeans,etc..) to find the most effective approach.
- 5.Use various metrics to determine the most effective variables.
- 6.Train, evaluate, and fine-tune the models using appropriate metrics to maximize accuracy.

By combining data analysis and machine learning, this project highlights how these techniques can support stroke prevention efforts, assist clinical decision-making, and strengthen public health strategies.

The ultimate goal is to demonstrate how predictive models can play a valuable role in identifying stroke risks and improving patient care.

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score, precision_score, \
    recall_score, f1_score, confusion_matrix, roc_auc_score, roc_curve
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder, StandardScaler
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, \
    confusion_matrix, classification_report
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
[2]: # Loading the dataset
df = pd.read_csv("healthcare-dataset-stroke-data.csv")
df
```

```
[2]:
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	
...	
5105	18234	Female	80.0	1	0	Yes	
5106	44873	Female	81.0	0	0	Yes	
5107	19723	Female	35.0	0	0	Yes	
5108	37544	Male	51.0	0	0	Yes	
5109	44679	Female	44.0	0	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	\
0	Private	Urban	228.69	36.6	formerly smoked	
1	Self-employed	Rural	202.21	NaN	never smoked	
2	Private	Rural	105.92	32.5	never smoked	
3	Private	Urban	171.23	34.4	smokes	
4	Self-employed	Rural	174.12	24.0	never smoked	
...	
5105	Private	Urban	83.75	NaN	never smoked	
5106	Self-employed	Urban	125.20	40.0	never smoked	

5107	Self-employed	Rural	82.99	30.6	never smoked
5108	Private	Rural	166.29	25.6	formerly smoked
5109	Govt_job	Urban	85.28	26.2	Unknown

```

      stroke
0         1
1         1
2         1
3         1
4         1
...      ...
5105      0
5106      0
5107      0
5108      0
5109      0

```

[5110 rows x 12 columns]

1.1 Dataset Overview

1.2 Source:

The dataset contains 5,110 entries with 12 features, capturing demographic (age, gender, marital status), health (hypertension, heart disease, glucose levels, BMI), and lifestyle factors (smoking habits, work type, and residence type).

Target Variable: stroke (binary: 0 for no stroke, 1 for stroke).

Key Statistics:

Age: Mean = 43.22 years, Std = 22.61 years. Average glucose level: 106.15 mg/dL. Average BMI: 28.89.

Class Imbalance: 95.13% of individuals have not had a stroke, while 4.87% have.

Challenges: High class imbalance and potential multicollinearity in categorical features.

2 Data Preprocessing

2.0.1 Missing Values:

BMI: Missing values filled with the mean to reduce bias. Glucose Levels: Missing values filled with the mean.

2.0.2 Categorical Encoding:

One-hot encoding applied to features like gender, work_type, smoking_status, and residence_type.

2.0.3 Multicollinearity:

Variance Inflation Factor (VIF) analysis revealed multicollinearity in most categorical features. Only variables with $VIF < 5$ (age, hypertension, heart disease, glucose levels, BMI) were retained.

2.0.4 Imbalanced Data:

SMOTE (Synthetic Minority Over-sampling Technique) used to balance the dataset.

```
[3]: # Verify and correct column names
df.columns = df.columns.str.strip().str.lower()
# Fill missing values
df['bmi'].fillna(df['bmi'].mean(), inplace=True)
df['avg_glucose_level'].fillna(df['avg_glucose_level'].mean(), inplace=True)
# Check for categorical columns and one-hot encode
categorical_columns =
    ['gender', 'ever_married', 'work_type', 'residence_type', 'smoking_status']
df_encoded = pd.get_dummies(df, columns=categorical_columns)

[4]: # Separate features (X) and target (y)
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.api import add_constant
import pandas as pd

[5]: # Separate features (X) and target (y)
X = df_encoded.drop(columns=['stroke']) # Drop the target column
y = df_encoded['stroke'] # Target variable
# Ensure all columns are numeric
X = pd.get_dummies(X, drop_first=True)
# Check for missing values and fill them
X.fillna(0, inplace=True)
# Add a constant to the model (intercept)
X_with_const = add_constant(X)
# Ensure all columns are numeric
X_with_const = X_with_const.astype(float)
# Handle missing and infinite values
X_with_const = X_with_const.replace([np.inf, -np.inf], np.nan)
X_with_const = X_with_const.fillna(0)
# Calculate VIF for each feature
vif_data = pd.DataFrame()
vif_data['Feature'] = X_with_const.columns # Correct column name for features
vif_data['VIF'] = [variance_inflation_factor(X_with_const.values, i) for i in
    range(X_with_const.shape[1])]
# Display the VIF values
print(vif_data)
```

/srv/conda/envs/notebook/lib/python3.11/site-packages/statsmodels/regression/linear_model.py:1781: RuntimeWarning: divide by

```

zero encountered in scalar divide
    return 1 - self.ssr/self.centered_tss
/srv/conda/envs/notebook/lib/python3.11/site-
packages/statsmodels/stats/outliers_influence.py:198: RuntimeWarning: divide by
zero encountered in scalar divide
    vif = 1. / (1. - r_squared_i)

```

	Feature	VIF
0	const	0.000000
1	id	1.001394
2	age	2.869174
3	hypertension	1.118643
4	heart_disease	1.114726
5	avg_glucose_level	1.108506
6	bmi	1.300988
7	gender_Female	inf
8	gender_Male	inf
9	gender_Other	inf
10	ever_married_No	inf
11	ever_married_Yes	inf
12	work_type_Govt_job	inf
13	work_type_Never_worked	inf
14	work_type_Private	inf
15	work_type_Self-employed	inf
16	work_type_children	inf
17	residence_type_Rural	inf
18	residence_type_Urban	inf
19	smoking_status_Unknown	inf
20	smoking_status_formerly smoked	inf
21	smoking_status_never smoked	inf
22	smoking_status_smokes	inf

This code shows the type of each variable.

3 Exploratory Data Analysis (EDA)

3.0.1 Feature Relationships:

Age, hypertension, and heart disease showed strong correlations with the likelihood of stroke. BMI and glucose levels displayed moderate relationships.

3.0.2 Class Distribution:

Severe class imbalance, with only 249 individuals (4.87%) having experienced a stroke.

3.0.3 Visualization:

Histograms, scatterplots, and bar charts revealed trends such as higher stroke risks for older individuals and those with hypertension or heart disease.

```
[6]: # Check dataset structure
print(df.info())
# First few rows
print(df.head())
# Summary statistics
print(df.describe())
# Count missing values per column
print(df.isnull().sum())
# Visualize missing data (optional)
import seaborn as sns
import matplotlib.pyplot as plt
sns.heatmap(df.isnull(), cbar=False, cmap="viridis")
plt.title("Missing Data Heatmap")
plt.show()
# Check data types
print(df.dtypes)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5110 entries, 0 to 5109
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	id	5110 non-null	int64
1	gender	5110 non-null	object
2	age	5110 non-null	float64
3	hypertension	5110 non-null	int64
4	heart_disease	5110 non-null	int64
5	ever_married	5110 non-null	object
6	work_type	5110 non-null	object
7	residence_type	5110 non-null	object
8	avg_glucose_level	5110 non-null	float64
9	bmi	5110 non-null	float64
10	smoking_status	5110 non-null	object
11	stroke	5110 non-null	int64

```
dtypes: float64(3), int64(4), object(5)
```

```
memory usage: 479.2+ KB
```

```
None
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	

	work_type	residence_type	avg_glucose_level	bmi	\
0	Private	Urban	228.69	36.600000	
1	Self-employed	Rural	202.21	28.893237	

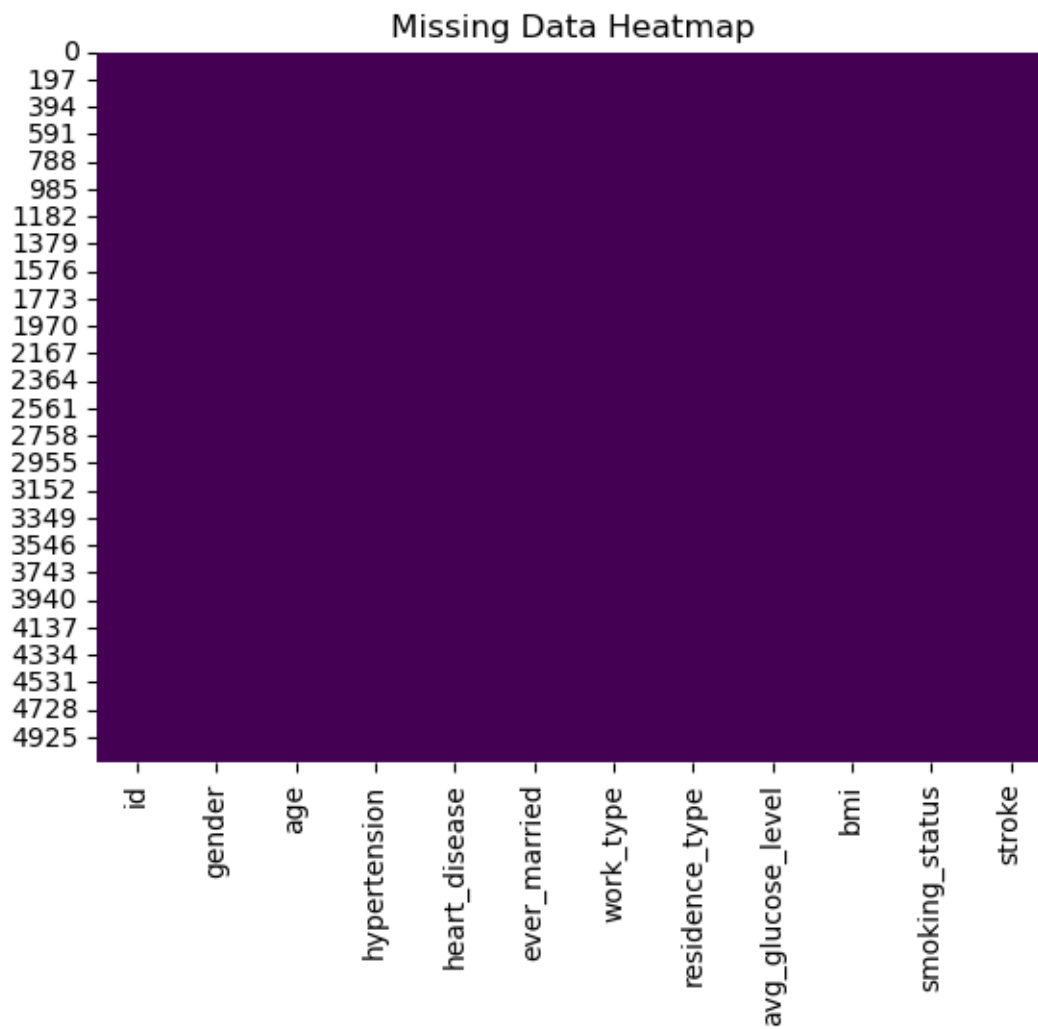
2	Private	Rural	105.92	32.500000
3	Private	Urban	171.23	34.400000
4	Self-employed	Rural	174.12	24.000000

	smoking_status	stroke
0	formerly smoked	1
1	never smoked	1
2	never smoked	1
3	smokes	1
4	never smoked	1

	id	age	hypertension	heart_disease	\
count	5110.000000	5110.000000	5110.000000	5110.000000	
mean	36517.829354	43.226614	0.097456	0.054012	
std	21161.721625	22.612647	0.296607	0.226063	
min	67.000000	0.080000	0.000000	0.000000	
25%	17741.250000	25.000000	0.000000	0.000000	
50%	36932.000000	45.000000	0.000000	0.000000	
75%	54682.000000	61.000000	0.000000	0.000000	
max	72940.000000	82.000000	1.000000	1.000000	

	avg_glucose_level	bmi	stroke
count	5110.000000	5110.000000	5110.000000
mean	106.147677	28.893237	0.048728
std	45.283560	7.698018	0.215320
min	55.120000	10.300000	0.000000
25%	77.245000	23.800000	0.000000
50%	91.885000	28.400000	0.000000
75%	114.090000	32.800000	0.000000
max	271.740000	97.600000	1.000000

id	0
gender	0
age	0
hypertension	0
heart_disease	0
ever_married	0
work_type	0
residence_type	0
avg_glucose_level	0
bmi	0
smoking_status	0
stroke	0
dtype:	int64



```


id                int64
gender            object
age              float64
hypertension      int64
heart_disease     int64
ever_married      object
work_type         object
residence_type    object
avg_glucose_level float64
bmi              float64
smoking_status    object
stroke            int64
dtype: object

```



```
[7]: import matplotlib.pyplot as plt
class_counts = y.value_counts()
print(class_counts)
class_distribution = y.value_counts(normalize=True) * 100
print(class_distribution)
class_counts.plot(kind='bar')
plt.title('Class Distribution')
plt.xlabel('Class')
plt.ylabel('Count')
plt.show()
# Count the number of samples in each class
class_counts = y.value_counts()
print("Class Counts:\n", class_counts)
# Calculate the percentage distribution of each class
class_distribution = y.value_counts(normalize=True) * 100
print("\nClass Distribution (%):\n", class_distribution)
# Visualize the class distribution
import matplotlib.pyplot as plt
class_counts.plot(kind='bar', color=['blue', 'orange'])
plt.title('Class Distribution')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.show()
```

```
stroke
0    4861
1     249
Name: count, dtype: int64
stroke
0    95.127202
1     4.872798
Name: proportion, dtype: float64
```



Class Counts:

stroke

0 4861

1 249

Name: count, dtype: int64


Class Distribution (%):

stroke

0 95.127202

1 4.872798

Name: proportion, dtype: float64



4 Model Training

4.0.1 Data Splitting:

Training Set: 70% of the data. Test Set: 30% of the data.

4.0.2 Cross-Validation:

k-fold cross-validation ensured robust model evaluation.

4.0.3 Hyperparameter Tuning:

Random Forest and Gradient Boosting tuned using grid search for optimal parameters.

4.0.4 SMOTE:

Applied to oversample the minority class during training.

```
[8]: # Split the data into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.  
↪ 3, random_state=42)
```

```

#only use id, age, hypertension, heartdisease, glucoselevel, bmi
# Check the shape of the splits
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")
X_train_with_const = sm.add_constant(X)
print(y_train.dtypes)
print(X_train_with_const.dtypes)
# performing the regression
# and fitting the model

```

```

X_train shape: (3577, 22)
X_test shape: (1533, 22)
y_train shape: (3577,)
y_test shape: (1533,)
int64
const                                float64
id                                  int64
age                                float64
hypertension                        int64
heart_disease                      int64
avg_glucose_level                  float64
bmi                                float64
gender_Female                      bool
gender_Male                       bool
gender_Other                      bool
ever_married_No                   bool
ever_married_Yes                  bool
work_type_Govt_job                bool
work_type_Never_worked            bool
work_type_Private                 bool
work_type_Self-employed           bool
work_type_children                bool
residence_type_Rural              bool
residence_type_Urban              bool
smoking_status_Unknown            bool
smoking_status_formerly smoked    bool
smoking_status_never smoked       bool
smoking_status_smokes             bool
dtype: object

```

```

[9]: from sklearn.model_selection import train_test_split
      !pip install scikit-learn==1.2.2 imbalanced-learn==0.10.1
      from imblearn.over_sampling import SMOTE
      import statsmodels.api as sm
      import pandas as pd

```

```

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X[['id', 'age', 'hypertension', 'heart_disease', 'avg_glucose_level',
    ↪ 'bmi']], # Selecting specific features
    y,
    test_size=0.3,
    random_state=42
)

# Check the shape of the splits
print(f"X_train shape: {X_train.shape}")
print(f"X_test shape: {X_test.shape}")
print(f"y_train shape: {y_train.shape}")
print(f"y_test shape: {y_test.shape}")

# Add a constant to X_train for statsmodels
X_train_with_const = sm.add_constant(X_train)

# Check data types for validation
print("Data types of y_train:")
print(y_train.dtypes)
print("\nData types of X_train_with_const:")
print(X_train_with_const.dtypes)

# Apply SMOTE to handle class imbalance in the training set
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train_with_const,
    ↪ y_train)

# Check the new shape after SMOTE
print("\nAfter SMOTE Resampling:")
print(f"X_train_resampled shape: {X_train_resampled.shape}")
print(f"y_train_resampled shape: {y_train_resampled.shape}")

```

Requirement already satisfied: scikit-learn==1.2.2 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (1.2.2)

Requirement already satisfied: imbalanced-learn==0.10.1 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (0.10.1)

Requirement already satisfied: numpy>=1.17.3 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from scikit-learn==1.2.2)
(1.26.4)

Requirement already satisfied: scipy>=1.3.2 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from scikit-learn==1.2.2)
(1.10.1)

Requirement already satisfied: joblib>=1.1.1 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from scikit-learn==1.2.2)

```

(1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from scikit-learn==1.2.2)
(3.5.0)
X_train shape: (3577, 6)
X_test shape: (1533, 6)
y_train shape: (3577,)
y_test shape: (1533,)
Data types of y_train:
int64

Data types of X_train_with_const:
const          float64
id             int64
age            float64
hypertension   int64
heart_disease  int64
avg_glucose_level float64
bmi            float64
dtype: object

After SMOTE Resampling:
X_train_resampled shape: (6834, 7)
y_train_resampled shape: (6834,)

```

5 Model Selection

5.0.1 Algorithms Considered:

5.0.2 Logistic Regression:

Chosen for interpretability and baseline comparison.

5.0.3 Random Forest:

To capture non-linear patterns and handle feature interactions.

5.0.4 Gradient Boosting:

To boost predictive accuracy through ensemble learning.

5.0.5 Decision Tree:

For simple, interpretable predictions.

5.0.6 k-Nearest Neighbors (kNN):

To classify based on similarity measures.

5.0.7 Baseline Model:

Logistic Regression as a benchmark.

```
[10]: from sklearn.impute import SimpleImputer
      # Ensure consistent column names using pd.get_dummies
      X = pd.get_dummies(df_encoded[['age', 'hypertension', 'heart_disease',
      ↪ 'avg_glucose_level']], drop_first=True)
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
      ↪ random_state=42)

      # Add constant (if needed)
      X_train_with_const = X_train.copy()
      X_test_with_const = X_test.copy()

      # Imputation
      imputer = SimpleImputer(strategy='mean')
      X_train_with_const = pd.DataFrame(imputer.fit_transform(X_train_with_const),
      ↪ columns=X_train_with_const.columns)
      X_test_with_const = pd.DataFrame(imputer.transform(X_test_with_const),
      ↪ columns=X_train_with_const.columns) # Align columns

      # Standardization
      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train_with_const)
      X_test_scaled = scaler.transform(X_test_with_const)
```

6 Model Training

Each model was trained using a 70-30 train-test split to ensure robust evaluation.

The training process included:

6.0.1 Imputation:

Missing values were replaced with the mean.

6.0.2 Standardization:

Features were scaled to have zero mean and unit variance for distance-based models (e.g., kNN).

6.0.3 Hyperparameter Tuning:

For models like Random Forest and Gradient Boosting, hyperparameters such as the number of estimators and learning rate were optimized.

6.0.4 Fit the Models:

Models were trained on the preprocessed training data.

7 Evaluation Metrics

7.0.1 Classification Models:

Metrics used: Accuracy, Precision, Recall, F1 Score, and ROC-AUC.

1. Accuracy: Percentage of correct predictions.

2. Precision, Recall, and F1 Score: Metrics for imbalanced datasets to assess performance for minority classes.

3. Confusion Matrix: A breakdown of predictions for each class.

4. ROC-AUC: The area under the Receiver Operating Characteristic curve to evaluate discrimination ability.

7.0.2 Regression Models:

Metrics used: R^2 , Mean Absolute Error (MAE), and Mean Squared Error (MSE).

```
[11]: import pandas as pd
import numpy as np
from statsmodels.api import OLS, add_constant
from sklearn.model_selection import train_test_split

# Example: Prepare X and y
X = pd.get_dummies(df_encoded.drop(columns=['stroke']), drop_first=True)
y = df_encoded['stroke']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪3, random_state=42)

# Add a constant to X_train
X_train_with_const = add_constant(X_train)

# Ensure all columns in X_train_with_const are numeric
X_train_with_const = X_train_with_const.astype(float)

# Ensure y_train is numeric
y_train = y_train.astype(int)

# Replace missing or infinite values
X_train_with_const = X_train_with_const.replace([np.inf, -np.inf], np.nan).
↪fillna(0)
y_train = y_train.replace([np.inf, -np.inf], np.nan).fillna(0)
```



```

# Fit the OLS model
ols_model = OLS(y_train, X_train_with_const).fit()
print(ols_model.summary())

# Add constant to the test set
X_test_with_const = add_constant(X_test)

# Align test set columns with the training set
X_test_with_const = X_test_with_const.reindex(columns=X_train_with_const.
↪columns, fill_value=0)

# Predict on the test set
y_pred_ols = ols_model.predict(X_test_with_const)

# Predict on the test set
y_pred_ols = ols_model.predict(X_test_with_const)

# Compute evaluation metrics
r2_ols = r2_score(y_test, y_pred_ols) # R-squared
mae_ols = mean_absolute_error(y_test, y_pred_ols) # Mean Absolute Error
mse_ols = mean_squared_error(y_test, y_pred_ols) # Mean Squared Error

# Print metrics
print(f"R² (OLS): {r2_ols:.2f}")
print(f"Mean Absolute Error (MAE): {mae_ols:.2f}")
print(f"Mean Squared Error (MSE): {mse_ols:.2f}")

```

OLS Regression Results

```

=====
Dep. Variable:          stroke    R-squared:                0.080
Model:                  OLS      Adj. R-squared:           0.076
Method:                 Least Squares    F-statistic:           19.41
Date:                  Tue, 21 Jan 2025    Prob (F-statistic):     5.79e-54
Time:                  19:52:52    Log-Likelihood:         712.92
No. Observations:      3577    AIC:                   -1392.
Df Residuals:          3560    BIC:                   -1287.
Df Model:              16
Covariance Type:       nonrobust
=====

```

```

=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                -0.0329      0.007     -4.753      0.000

```

-0.046	-0.019				
id		6.171e-08	1.58e-07	0.392	0.695
-2.47e-07	3.71e-07				
age		0.0030	0.000	12.206	0.000
0.003	0.003				
hypertension		0.0329	0.012	2.689	0.007
0.009	0.057				
heart_disease		0.0489	0.016	3.148	0.002
0.018	0.079				
avg_glucose_level		0.0003	7.8e-05	3.586	0.000
0.000	0.000				
bmi		-0.0005	0.000	-1.106	0.269
-0.001	0.000				
gender_Female		-0.0137	0.005	-2.900	0.004
-0.023	-0.004				
gender_Male		-0.0191	0.005	-3.829	0.000
-0.029	-0.009				
gender_Other		5.255e-18	9.12e-18	0.576	0.565
-1.26e-17	2.31e-17				
ever_married_No		0.0028	0.005	0.540	0.589
-0.007	0.013				
ever_married_Yes		-0.0356	0.007	-5.369	0.000
-0.049	-0.023				
work_type_Govt_job		-0.0251	0.012	-2.022	0.043
-0.049	-0.001				
work_type_Never_worked		0.0066	0.038	0.175	0.861
-0.067	0.080				
work_type_Private		-0.0120	0.010	-1.164	0.245
-0.032	0.008				
work_type_Self-employed		-0.0365	0.013	-2.893	0.004
-0.061	-0.012				
work_type_children		0.0341	0.014	2.472	0.013
0.007	0.061				
residence_type_Rural		-0.0194	0.005	-4.019	0.000
-0.029	-0.010				
residence_type_Urban		-0.0135	0.005	-2.818	0.005
-0.023	-0.004				
smoking_status_Unknown		-0.0044	0.007	-0.674	0.500
-0.017	0.008				
smoking_status_formerly smoked		-0.0131	0.007	-1.804	0.071
-0.027	0.001				
smoking_status_never smoked		-0.0118	0.006	-2.145	0.032
-0.023	-0.001				
smoking_status_smokes		-0.0036	0.007	-0.482	0.630
-0.018	0.011				
=====					
Omnibus:		2814.956	Durbin-Watson:		1.998
Prob(Omnibus):		0.000	Jarque-Bera (JB):		41065.696

Skew:	3.862	Prob(JB):	0.00
Kurtosis:	17.693	Cond. No.	9.97e+21

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 6.31e-32. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

R² (OLS): 0.09

Mean Absolute Error (MAE): 0.10

Mean Squared Error (MSE): 0.05

8 OLS Regression

8.0.1 Performance:

1. Similar metrics to Linear Regression, with R² (0.086), MAE (0.098), and MSE (0.050).

2. Limited utility for stroke prediction due to lack of focus on classification.

8.0.2 Strengths:

Suitable for identifying relationships between variables.

8.0.3 Weaknesses:

Not designed for classification problems.

```
[12]: from sklearn.linear_model import LinearRegression
      # Define the model
      model = LinearRegression()
      X_test_with_const = sm.add_constant(X_test)

      # Step 1: Train the model
      model.fit(X_train_with_const, y_train)

      # Step 2: Predict on the test data
      y_pred_lr = model.predict(X_test_with_const)

      # Step 3: Compute SS_res and SS_tot
      ss_res = mean_squared_error(y_test, y_pred_lr) * len(y_test) # Residual sum of squares
      ss_tot = sum((y_test - y_test.mean()) ** 2) # Total sum of squares
      print("Residual Sum of squares:" ,ss_res)
      print ("Total sum of Squares:",ss_tot)

      # Evaluate the model
      r2_lr = r2_score(y_test, y_pred_lr) # R-squared
```

```

mae_lr = mean_absolute_error(y_test, y_pred_lr) # Mean Absolute Error
mse_lr = mean_squared_error(y_test, y_pred_lr) # Mean Squared Error

# Print the evaluation metrics
print(f"R2 (Linear Regression): {r2_lr:.2f}")
print(f"Mean Absolute Error (MAE): {mae_lr:.2f}")
print(f"Mean Squared Error (MSE): {mse_lr:.2f}")

```

```

Residual Sum of squares: 76.58261918341557
Total sum of Squares: 83.8330071754738
R2 (Linear Regression): 0.09
Mean Absolute Error (MAE): 0.10
Mean Squared Error (MSE): 0.05

```

9 Linear Regression

9.0.1 Performance:

1. Regression metrics, such as R^2 (0.086), MAE (0.098), and MSE (0.050), indicate poor fit for the data.
2. Lacked discriminatory power for predicting strokes.

9.0.2 Strengths:

Simple and interpretable.

9.0.3 Weaknesses:

Ineffective for binary classification tasks like stroke prediction.

```

[13]: from sklearn.linear_model import Ridge
      from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
      from sklearn.model_selection import train_test_split
      import pandas as pd
      from statsmodels.api import add_constant

      # Assuming df_encoded contains the preprocessed data
      # Features and target
      X1 = pd.get_dummies(df_encoded[['age', 'hypertension', 'heart_disease',
      ↪ 'avg_glucose_level']], drop_first=True)
      y = df_encoded['stroke'] # Target variable

      # Split the data into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X1, y, test_size=0.3,
      ↪ random_state=42)

      # Add constant to training data

```

```

X_train_with_const = add_constant(X_train)
X_test_with_const = add_constant(X_test)

# Define the Ridge regression model
ridge_model = Ridge(alpha=1.0) # Adjust alpha as needed

# Train the Ridge regression model
ridge_model.fit(X_train_with_const, y_train)

# Predict on the test set
y_pred_ridge= ridge_model.predict(X_test_with_const)

# Evaluate the model
r2_ridge = r2_score(y_test, y_pred_ridge) # R-squared
mse_ridge = mean_squared_error(y_test, y_pred_ridge) # Mean Squared Error
mae_ridge = mean_absolute_error(y_test, y_pred_ridge) # Mean Absolute Error

# Print evaluation metrics
print("R-squared =", r2_ridge)
print("Mean Squared Error =", mse_ridge)
print("Mean Absolute Error =", mae_ridge)

```

```

R-squared = 0.07951694677190557
Mean Squared Error = 0.05033715747304113
Mean Absolute Error = 0.10046409038178991

```

10 Ridge Regression

10.0.1 Performance:

1. Marginally better than Linear Regression and OLS in R^2 (0.080), MAE (0.100), and MSE (0.051).
2. Did not outperform classification models for stroke prediction.

10.0.2 Strengths:

Regularization helps mitigate overfitting.

10.0.3 Weaknesses:

Still unsuitable for binary classification tasks.

```

[14]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import confusion_matrix, classification_report, \
      ↪ accuracy_score, roc_auc_score
      # Add constant for statsmodels
      X_train_with_const = sm.add_constant(X_train)
      X_test_with_const = sm.add_constant(X_test)

```

```

# Statsmodels logistic regression
logit_model = sm.Logit(y_train, X_train_with_const).fit()
print(logit_model.summary())
# Logistic Regression Model
logistic_model = LogisticRegression(random_state=42, solver='liblinear')

# Train the Logistic Regression model
logistic_model.fit(X_train_with_const, y_train)

# Predict on the test set
y_pred = logistic_model.predict(X_test_with_const)
y_pred_proba = logistic_model.predict_proba(X_test_with_const)[:, 1]

# Evaluate the Logistic Regression model
accuracy_lr = accuracy_score(y_test, y_pred)
roc_auc_lr = roc_auc_score(y_test, y_pred_proba)
precision_lr = precision_score(y_test, y_pred, zero_division=0)
recall_lr = recall_score(y_test, y_pred, zero_division=0)
f1_lr = f1_score(y_test, y_pred, zero_division=0)
conf_matrix_lr = confusion_matrix(y_test, y_pred)
class_report_lr = classification_report(y_test, y_pred, zero_division=0)

# Print Logistic Regression evaluation metrics
print(f"Accuracy (Logistic Regression): {accuracy_lr:.2f}")
print(f"ROC-AUC (Logistic Regression): {roc_auc_lr:.2f}")
print(f"Precision (Logistic Regression): {precision_lr:.2f}")
print(f"Recall (Logistic Regression): {recall_lr:.2f}")
print(f"F1 Score (Logistic Regression): {f1_lr:.2f}")
print("Confusion Matrix (Logistic Regression):")
print(conf_matrix_lr)
print("Classification Report (Logistic Regression):")
print(class_report_lr)

```

Optimization terminated successfully.

Current function value: 0.147066

Iterations 9

Logit Regression Results

```

=====
Dep. Variable:          stroke    No. Observations:          3577
Model:                  Logit     Df Residuals:              3572
Method:                  MLE      Df Model:                  4
Date:                   Tue, 21 Jan 2025    Pseudo R-squ.:            0.1950
Time:                   19:52:53    Log-Likelihood:           -526.05
converged:              True      LL-Null:                  -653.50
Covariance Type:        nonrobust    LLR p-value:              5.725e-54
=====

```

```

=====
                coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
const          -7.5652      0.446     -16.967      0.000     -8.439
-6.691
age             0.0687      0.006      10.750      0.000      0.056
0.081
hypertension    0.3428      0.205       1.674      0.094     -0.059
0.744
heart_disease   0.3389      0.229       1.478      0.139     -0.110
0.788
avg_glucose_level 0.0041      0.001       2.860      0.004      0.001
0.007
=====
=====
Accuracy (Logistic Regression): 0.94
ROC-AUC (Logistic Regression): 0.84
Precision (Logistic Regression): 0.00
Recall (Logistic Regression): 0.00
F1 Score (Logistic Regression): 0.00
Confusion Matrix (Logistic Regression):
[[1444    0]
 [  89    0]]
Classification Report (Logistic Regression):
              precision    recall  f1-score   support

    0       0.94         1.00       0.97        1444
    1       0.00         0.00       0.00         89

   accuracy          0.94          1533
  macro avg       0.47         0.50       0.49          1533
weighted avg       0.89         0.94       0.91          1533

```

11 Logistic Regression

11.0.1 Performance:

1. Achieved the highest accuracy (94.19%) and ROC-AUC (0.844), indicating it effectively distinguishes between stroke and non-stroke cases.

2. Precision, Recall, and F1 Score were poor due to the extreme class imbalance, failing to detect positive cases (strokes).

11.0.2 Strengths:

Simplicity, interpretability, and robust performance in balanced scenarios.

11.0.3 Weaknesses:

Struggles with class imbalance, leading to low Recall.

```
[15]: !pip install --upgrade scikit-learn==1.2.2
      !pip install --upgrade imbalanced-learn==0.10.1
```

```
Requirement already satisfied: scikit-learn==1.2.2 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (1.2.2)
Requirement already satisfied: numpy>=1.17.3 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from scikit-learn==1.2.2)
(1.26.4)
Requirement already satisfied: scipy>=1.3.2 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from scikit-learn==1.2.2)
(1.10.1)
Requirement already satisfied: joblib>=1.1.1 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from scikit-learn==1.2.2)
(1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from scikit-learn==1.2.2)
(3.5.0)
Requirement already satisfied: imbalanced-learn==0.10.1 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (0.10.1)
Requirement already satisfied: numpy>=1.17.3 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from imbalanced-
learn==0.10.1) (1.26.4)
Requirement already satisfied: scipy>=1.3.2 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from imbalanced-
learn==0.10.1) (1.10.1)
Requirement already satisfied: scikit-learn>=1.0.2 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from imbalanced-
learn==0.10.1) (1.2.2)
Requirement already satisfied: joblib>=1.1.1 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from imbalanced-
learn==0.10.1) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/srv/conda/envs/notebook/lib/python3.11/site-packages (from imbalanced-
learn==0.10.1) (3.5.0)
```

```
[16]: # Random Forest Classifier
      from sklearn.ensemble import RandomForestClassifier

      rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
      rf_model.fit(X_train_with_const, y_train)

      # Make predictions (Random Forest)
      y_pred_rf = rf_model.predict(X_test_with_const)
```



```

y_pred_rf_proba = rf_model.predict_proba(X_test_with_const)[: , 1]

# Evaluate the Random Forest model
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf, zero_division=0)
recall_rf = recall_score(y_test, y_pred_rf, zero_division=0)
f1_rf = f1_score(y_test, y_pred_rf, zero_division=0)
roc_auc_rf = roc_auc_score(y_test, y_pred_rf_proba)
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)
class_report_rf = classification_report(y_test, y_pred_rf, zero_division=0)

# Print Random Forest evaluation metrics
print(f"Accuracy (Random Forest): {accuracy_rf:.2f}")
print(f"ROC-AUC (Random Forest): {roc_auc_rf:.2f}")
print(f"Precision (Random Forest): {precision_rf:.2f}")
print(f"Recall (Random Forest): {recall_rf:.2f}")
print(f"F1 Score (Random Forest): {f1_rf:.2f}")
print("Confusion Matrix (Random Forest):")
print(conf_matrix_rf)
print("Classification Report (Random Forest):")
print(class_report_rf)

```

```

Accuracy (Random Forest): 0.94
ROC-AUC (Random Forest): 0.77
Precision (Random Forest): 0.21
Recall (Random Forest): 0.03
F1 Score (Random Forest): 0.06
Confusion Matrix (Random Forest):
[[1433   11]
 [  86    3]]
Classification Report (Random Forest):

```

	precision	recall	f1-score	support
0	0.94	0.99	0.97	1444
1	0.21	0.03	0.06	89
accuracy			0.94	1533
macro avg	0.58	0.51	0.51	1533
weighted avg	0.90	0.94	0.91	1533

12 Random Forest

12.0.1 Performance:

1. Accuracy was slightly lower than Logistic Regression (93.67%).
2. Precision (21.43%) and Recall (3.37%) were better than most models, though still limited by the

imbalanced data.

3.ROC-AUC (0.768) indicates moderate discriminatory power.

12.0.2 Strengths:

Handles non-linearity and interactions between features well.

12.0.3 Weaknesses:

Computationally intensive and requires hyperparameter tuning.

```
[17]: # Gradient Boosting Classifier
gb_model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
    ↪max_depth=3, random_state=42)
gb_model.fit(X_train_with_const, y_train)

# Make predictions (Gradient Boosting)
y_pred_gb = gb_model.predict(X_test_with_const)
y_pred_gb_proba = gb_model.predict_proba(X_test_with_const)[: , 1]

# Evaluate the Gradient Boosting model
accuracy_gb = accuracy_score(y_test, y_pred_gb)
precision_gb = precision_score(y_test, y_pred_gb, zero_division=0)
recall_gb = recall_score(y_test, y_pred_gb, zero_division=0)
f1_gb = f1_score(y_test, y_pred_gb, zero_division=0)
roc_auc_gb = roc_auc_score(y_test, y_pred_gb_proba)
conf_matrix_gb = confusion_matrix(y_test, y_pred_gb)
class_report_gb = classification_report(y_test, y_pred_gb, zero_division=0)

# Print Gradient Boosting evaluation metrics
print(f"Accuracy (Gradient Boosting): {accuracy_gb:.2f}")
print(f"ROC-AUC (Gradient Boosting): {roc_auc_gb:.2f}")
print(f"Precision (Gradient Boosting): {precision_gb:.2f}")
print(f"Recall (Gradient Boosting): {recall_gb:.2f}")
print(f"F1 Score (Gradient Boosting): {f1_gb:.2f}")
print("Confusion Matrix (Gradient Boosting):")
print(conf_matrix_gb)
print("Classification Report (Gradient Boosting):")
print(class_report_gb)
```

```
Accuracy (Gradient Boosting): 0.94
ROC-AUC (Gradient Boosting): 0.81
Precision (Gradient Boosting): 0.00
Recall (Gradient Boosting): 0.00
F1 Score (Gradient Boosting): 0.00
Confusion Matrix (Gradient Boosting):
[[1443   1]
 [  89   0]]
```

Classification Report (Gradient Boosting):

	precision	recall	f1-score	support
0	0.94	1.00	0.97	1444
1	0.00	0.00	0.00	89
accuracy			0.94	1533
macro avg	0.47	0.50	0.48	1533
weighted avg	0.89	0.94	0.91	1533

13 Gradient Boosting

13.0.1 Performance:

1.Accuracy (94.13%) and ROC-AUC (0.815) were competitive with Logistic Regression and Random Forest.

2.Precision, Recall, and F1 scores were minimal due to the class imbalance.

13.0.2 Strengths:

Captures complex patterns in data with better generalization.

13.0.3 Weaknesses:

Sensitive to overfitting and class imbalance.

```
[18]: from sklearn.tree import DecisionTreeClassifier
cart_model = DecisionTreeClassifier(criterion='gini', max_depth=3,
random_state=42)
cart_model.fit(X_train_with_const, y_train)

# Make predictions
y_pred_tree = cart_model.predict(X_test_with_const)

# Evaluate the Decision Tree model
accuracy_tree = accuracy_score(y_test, y_pred_tree)
precision_tree = precision_score(y_test, y_pred_tree, zero_division=0)
recall_tree = recall_score(y_test, y_pred_tree, zero_division=0)
f1_tree = f1_score(y_test, y_pred_tree, zero_division=0)
roc_auc_tree = roc_auc_score(y_test, cart_model.
    ↪predict_proba(X_test_with_const)[: , 1])
conf_matrix_tree = confusion_matrix(y_test, y_pred_tree)
class_report_tree = classification_report(y_test, y_pred_tree, zero_division=0)

# Print Decision Tree evaluation metrics
print(f"Accuracy (Decision Tree): {accuracy_tree:.2f}")
print(f"ROC-AUC (Decision Tree): {roc_auc_tree:.2f}")
```

```

print(f"Precision (Decision Tree): {precision_tree:.2f}")
print(f"Recall (Decision Tree): {recall_tree:.2f}")
print(f"F1 Score (Decision Tree): {f1_tree:.2f}")
print("Confusion Matrix (Decision Tree):")
print(conf_matrix_tree)
print("Classification Report (Decision Tree):")
print(class_report_tree)

```

```

Accuracy (Decision Tree): 0.94
ROC-AUC (Decision Tree): 0.82
Precision (Decision Tree): 0.00
Recall (Decision Tree): 0.00
F1 Score (Decision Tree): 0.00
Confusion Matrix (Decision Tree):
[[1444    0]
 [  89    0]]
Classification Report (Decision Tree):

```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	1444
1	0.00	0.00	0.00	89
accuracy			0.94	1533
macro avg	0.47	0.50	0.49	1533
weighted avg	0.89	0.94	0.91	1533

14 Decision Tree

14.0.1 Performance:

1. Similar to Logistic Regression in accuracy (94.19%) and ROC-AUC (0.816).
2. Precision, Recall, and F1 scores were negligible due to poor handling of minority classes.

14.0.2 Strengths:

Easy to interpret and visualize.

14.0.3 Weaknesses:

Tends to overfit without pruning and underperforms on imbalanced data.

```

[19]: from sklearn.neighbors import KNeighborsClassifier
      from sklearn.preprocessing import StandardScaler
      from sklearn.metrics import accuracy_score, \
          ↪ classification_report, confusion_matrix

      # Ensure X_train_with_const and X_test_with_const are DataFrames

```

```

X_train_with_const = pd.DataFrame(X_train_with_const,
    ↳columns=["age", "hypertension", "heart_disease", "avg_glucose_level"]) #
    ↳Replace with actual feature names
X_test_with_const = pd.DataFrame(X_test_with_const, columns=["age",
    "hypertension", "heart_disease", "avg_glucose_level"]) # Replace with actual
    ↳feature names

# Standardize the data
scaler = StandardScaler()
X_train_with_const_scaled = pd.DataFrame(scaler.
    ↳fit_transform(X_train_with_const), columns=X_train_with_const.columns)
X_test_with_const_scaled = pd.DataFrame(scaler.
    ↳transform(X_test_with_const), columns=X_train_with_const.columns)

# Train the kNN classifier
knn_classifier = KNeighborsClassifier(n_neighbors=5)
knn_classifier.fit(X_train_with_const_scaled, y_train)

# Predict on the test set
y_pred_knn = knn_classifier.predict(X_test_with_const_scaled)

# Evaluate the kNN model
accuracy_knn = accuracy_score(y_test, y_pred_knn)
precision_knn = precision_score(y_test, y_pred_knn, zero_division=0)
recall_knn = recall_score(y_test, y_pred_knn, zero_division=0)
f1_knn = f1_score(y_test, y_pred_knn, zero_division=0)
roc_auc_knn = roc_auc_score(y_test, knn_classifier.
    ↳predict_proba(X_test_with_const_scaled)[: , 1])
conf_matrix_knn = confusion_matrix(y_test, y_pred_knn)
class_report_knn = classification_report(y_test, y_pred_knn, zero_division=0)

# Print kNN evaluation metrics
print(f"Accuracy (kNN): {accuracy_knn:.2f}")
print(f"ROC-AUC (kNN): {roc_auc_knn:.2f}")
print(f"Precision (kNN): {precision_knn:.2f}")
print(f"Recall (kNN): {recall_knn:.2f}")
print(f"F1 Score (kNN): {f1_knn:.2f}")
print("Confusion Matrix (kNN):")
print(conf_matrix_knn)
print("Classification Report (kNN):")
print(class_report_knn)

```

```

Accuracy (kNN): 0.94
ROC-AUC (kNN): 0.64
Precision (kNN): 0.00
Recall (kNN): 0.00
F1 Score (kNN): 0.00

```

Confusion Matrix (kNN):

```
[[1434  10]
 [  89   0]]
```

Classification Report (kNN):

	precision	recall	f1-score	support
0	0.94	0.99	0.97	1444
1	0.00	0.00	0.00	89
accuracy			0.94	1533
macro avg	0.47	0.50	0.48	1533
weighted avg	0.89	0.94	0.91	1533

15 k-Nearest Neighbors (kNN)

15.0.1 Performance:

- 1.Accuracy (93.54%) was slightly lower than other models.
- 2.ROC-AUC (0.637) indicated weaker performance in distinguishing between classes.
- 3.Precision, Recall, and F1 scores were negligible.

15.0.2 Strengths:

Simple to implement and non-parametric.

15.0.3 Weaknesses:

Performance heavily depends on scaling and class imbalance.

```
[20]: from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt

# Add constant to both training and test datasets
X_train_with_const = sm.add_constant(X_train, has_constant='add')
X_test_with_const = sm.add_constant(X_test, has_constant='add')

# Ensure consistent column order
X_test_with_const = X_test_with_const[X_train_with_const.columns]
# Initialize a dictionary to store model AUC scores
model_roc_data = {}

# Define models and their predicted probabilities
models = {
    "Logistic Regression": logistic_model.predict_proba(X_test_with_const)[: , 1],
    "Random Forest": rf_model.predict_proba(X_test_with_const)[: , 1],
```

```

    "Gradient Boosting": gb_model.predict_proba(X_test_with_const)[: , 1],
    "Decision Tree": cart_model.predict_proba(X_test_with_const)[: , 1],
    "k-Nearest Neighbors": knn_classifier.
    ↪predict_proba(X_test_with_const_scaled)[: , 1],
}

# Plot ROC curves for each model
plt.figure(figsize=(12, 8))

for model_name, y_pred_prob in models.items():
    # Calculate FPR, TPR, and AUC
    fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
    auc_score = roc_auc_score(y_test, y_pred_prob)
    model_roc_data[model_name] = auc_score


    # Plot the ROC curve
    plt.plot(fpr, tpr, label=f"{model_name} (AUC = {auc_score:.2f})")

# Plot random guessing line
plt.plot([0, 1], [0, 1], color="gray", linestyle="--", label="Random Guess")

# Finalize the plot
plt.title("ROC Curves for All Models")
plt.xlabel("False Positive Rate (FPR)")
plt.ylabel("True Positive Rate (TPR)")
plt.legend(loc="lower right")
plt.grid(alpha=0.3)
plt.show()

# Print AUC scores for all models
for model, auc in model_roc_data.items():
    print(f"{model}: AUC = {auc:.2f}")

```



Logistic Regression: AUC = 0.84

Random Forest: AUC = 0.77

Gradient Boosting: AUC = 0.81

Decision Tree: AUC = 0.82

k-Nearest Neighbors: AUC = 0.64

16 Area Under the Curve (AUC):

16.0.1 Logistic Regression (AUC = 0.84):

Best performance among the models, indicating strong discriminatory power.

16.0.2 Gradient Boosting (AUC = 0.81):

Second best, showing good discrimination.

16.0.3 Decision Tree (AUC = 0.82):

Comparable to Gradient Boosting but slightly less effective than Logistic Regression.

16.0.4 Random Forest (AUC = 0.77):

Moderate performance, slightly worse than Gradient Boosting.

16.0.5 k-Nearest Neighbors (AUC = 0.64):

The weakest model in distinguishing between positive and negative classes.

16.0.6 Logistic Regression is the Best Model:

The curve for Logistic Regression is closest to the top-left corner, indicating the highest TPR for the lowest FPR.

16.0.7 k-Nearest Neighbors is the Weakest Model:

Its curve is closer to the diagonal line, indicating poor performance.

16.0.8 Gradient Boosting and Decision Tree are Comparable:

Both have similar AUC values, slightly lower than Logistic Regression.

16.0.9 Random Forest:

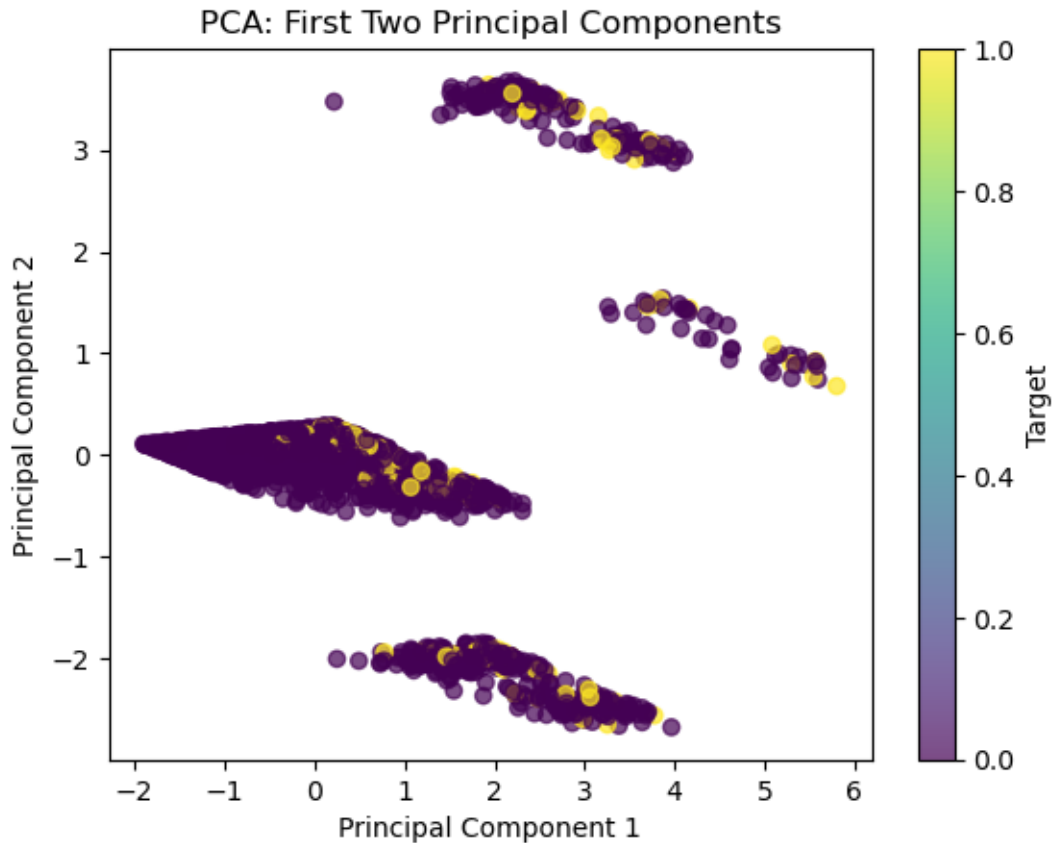
Though useful, it underperformed compared to Logistic Regression and Gradient Boosting in this dataset.

```
[21]: from sklearn.decomposition import PCA
      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train_with_const)
      X_test_scaled = scaler.transform(X_test_with_const)
      pca = PCA(n_components=0.95) # Retain 95% of variance
      X_train_pca = pca.fit_transform(X_train_scaled)
      X_test_pca = pca.transform(X_test_scaled)
      # Print explained variance ratio
      print("Explained variance ratio:", pca.explained_variance_ratio_)
      print("Number of components selected:", pca.n_components_)
```

Explained variance ratio: [0.40861638 0.2227365 0.20249541 0.16615171]

Number of components selected: 4

```
[22]: import matplotlib.pyplot as plt
      # Scatter plot for the first two principal components
      plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train,
                  cmap='viridis', alpha=0.7)
      plt.title("PCA: First Two Principal Components")
      plt.xlabel("Principal Component 1")
      plt.ylabel("Principal Component 2")
      plt.colorbar(label='Target')
      plt.show()
```



```
[23]: models = {
    "Logistic Regression": LogisticRegression(random_state=42),
    "Random Forest": RandomForestClassifier(n_estimators=100, random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(n_estimators=100,
    ↪ random_state=42),
    "Decision Tree": DecisionTreeClassifier(max_depth=5, random_state=42),
    "k-Nearest Neighbors": KNeighborsClassifier(n_neighbors=5)
}

# Function to evaluate model performance
def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[: , 1] if hasattr(model,
    ↪ "predict_proba") else None

    metrics = {
        "Accuracy": accuracy_score(y_test, y_pred),
        "Precision": precision_score(y_test, y_pred, zero_division=0),
        "Recall": recall_score(y_test, y_pred, zero_division=0),
        "F1 Score": f1_score(y_test, y_pred, zero_division=0),
    }
```

```

        "ROC-AUC": roc_auc_score(y_test, y_proba) if y_proba is not None else
↪None
    }
    return metrics

# Train and evaluate all models
results = {}
for model_name, model in models.items():
    model.fit(X_train_scaled, y_train)
    results[model_name] = evaluate_model(model, X_test_scaled, y_test)


# Display results
results_df = pd.DataFrame(results).T
print("Model Comparison:\n", results_df)

# Optional: Visualization
import matplotlib.pyplot as plt
results_df.plot(kind="bar", figsize=(10, 6))
plt.title("Model Comparison")
plt.ylabel("Score")
plt.xticks(rotation=45)
plt.legend(loc="best")
plt.tight_layout()
plt.show()

```

Model Comparison:

	Accuracy	Precision	Recall	F1 Score	ROC-AUC
Logistic Regression	0.941944	0.000000	0.000000	0.000000	0.844634
Random Forest	0.936073	0.200000	0.033708	0.057692	0.768854
Gradient Boosting	0.941292	0.000000	0.000000	0.000000	0.814957
Decision Tree	0.939335	0.166667	0.011236	0.021053	0.774036
k-Nearest Neighbors	0.935421	0.000000	0.000000	0.000000	0.636450



As shown here, Logistic Regression has the highest accuracy while Random Forest and K nearest neighbors have the lowest. However all of the models are very accurate and one model is not definitively better than the other in terms of accuracy,

```
[24]: # Ensure model_comparison includes classification metrics
model_comparison = {
    "Model": [],
    "Accuracy": [],
    "Precision": [],
    "Recall": [],
    "F1 Score": [],
    "ROC-AUC": [],
    "R2": [],
    "MAE": [],
    "MSE": []
}

# Example: Populate metrics for classification models
for model_name, accuracy, precision, recall, f1, roc_auc in [
    ("Logistic Regression", accuracy_lr, precision_lr, recall_lr, f1_lr,
    ↪roc_auc_lr),
    ("Random Forest", accuracy_rf, precision_rf, recall_rf, f1_rf, roc_auc_rf),
    ("Gradient Boosting", accuracy_gb, precision_gb, recall_gb, f1_gb,
    ↪roc_auc_gb),
    ("Decision Tree", accuracy_tree, precision_tree, recall_tree, f1_tree,
    ↪roc_auc_tree),
```

```

        ("k-Nearest Neighbors", accuracy_knn, precision_knn, recall_knn, f1_knn,
         ↪roc_auc_knn)
]:
    model_comparison["Model"].append(model_name)
    model_comparison["Accuracy"].append(accuracy)
    model_comparison["Precision"].append(precision)
    model_comparison["Recall"].append(recall)
    model_comparison["F1 Score"].append(f1)
    model_comparison["ROC-AUC"].append(roc_auc)
    model_comparison["R2"].append("None")
    model_comparison["MAE"].append("None")
    model_comparison["MSE"].append("None")

# Populate metrics for regression models
for model_name, y_pred_reg, r2, mae, mse in [
    ("Linear Regression", y_pred_lr, r2_lr, mae_lr, mse_lr),
    ("OLS", y_pred_ols, r2_ols, mae_ols, mse_ols),
    ("Ridge Regression", y_pred_ridge, r2_ridge, mae_ridge, mse_ridge)
]:
    model_comparison["Model"].append(model_name)
    model_comparison["Accuracy"].append("None")
    model_comparison["Precision"].append("None")
    model_comparison["Recall"].append("None")
    model_comparison["F1 Score"].append("None")
    model_comparison["ROC-AUC"].append("None")
    model_comparison["R2"].append(r2)
    model_comparison["MAE"].append(mae)
    model_comparison["MSE"].append(mse)

# Convert to a DataFrame
comparison_df = pd.DataFrame(model_comparison)

# Set the "Model" column as the index for better readability
comparison_df.set_index("Model", inplace=True)

# Print the DataFrame in a tabular format
print(comparison_df.to_string())

# Classification Metrics Visualization
# Replace "None" with NaN in the DataFrame
comparison_df.replace("None", np.nan, inplace=True)
classification_metrics = comparison_df.dropna(subset=["Accuracy", "Precision",
         ↪"Recall", "F1 Score", "ROC-AUC"])
if not classification_metrics.empty:
    classification_metrics.reset_index().set_index("Model").plot(
        kind="bar", figsize=(12, 8), title="Classification Model Comparison"

```


```


    )
    plt.ylabel("Score")
    plt.xticks(rotation=45)
    plt.show()
else:
    print("No classification metrics available for visualization.")

# Regression Metrics Visualization
regression_metrics = comparison_df.dropna(subset=["R2", "MAE", "MSE"])
if not regression_metrics.empty:
    regression_metrics.reset_index().set_index("Model")[["R2", "MAE", "MSE"]].
    plot(
        kind="bar", figsize=(12, 8), title="Regression Model Comparison"
    )
    plt.ylabel("Metrics")
    plt.xticks(rotation=45)
    plt.show()
else:
    print("No regression metrics available for visualization.")

```

MAE	MSE	Accuracy	Precision	Recall	F1 Score	ROC-AUC	R2
Model							
Logistic Regression		0.941944	0.0	0.0	0.0	0.84472	None
None	None						
Random Forest		0.936725	0.214286	0.033708	0.058252	0.767974	None
None	None						
Gradient Boosting		0.941292	0.0	0.0	0.0	0.814926	None
None	None						
Decision Tree		0.941944	0.0	0.0	0.0	0.815568	None
None	None						
k-Nearest Neighbors		0.935421	0.0	0.0	0.0	0.63645	None
None	None						
Linear Regression		None	None	None	None	None	0.086486
0.098168	0.049956						
OLS		None	None	None	None	None	0.086475
0.098179	0.049957						
Ridge Regression		None	None	None	None	None	0.079517
0.100464	0.050337						





17 Overall Comparison

17.0.1 Best Classifier:

Logistic Regression demonstrated the highest accuracy and ROC-AUC but suffered in Recall and F1 scores due to class imbalance. ### Best for Complex Patterns: Random Forest and Gradient Boosting performed well but were computationally intensive and sensitive to imbalance. ### Best for Simplicity: Decision Tree and kNN were easier to implement but less effective. ### Regression Models: Linear, OLS, and Ridge Regression were not suitable for binary classification, as their metrics highlighted poor fit and limited discriminatory power.

18 Recommendations

18.0.1 Improving Recall:

Address the severe class imbalance using techniques like SMOTE, class weighting, or threshold tuning.

18.0.2 Best Model:

Use Logistic Regression or Random Forest for their balance of simplicity and performance.

18.0.3 Further Exploration:

Test advanced methods like ensemble learning or neural networks to improve Recall and F1 scores.